

Measuring breakthrough-ness on new papers — how far a GRO metric gets, and the ceiling above it

Written 2026-07-09 (v2, supersedes the 2026-07-08 draft). Companion to `breakthrough-metric.md` ("Do we have a scientific-breakthrough metric? — assessment + design"), which argued no metric in the program measures breakthroughs, decomposed "breakthrough" into three signals (disruption / significance / realized uptake), and named three blockers.

Abstract

We asked whether a metric computed over GRO's structured sidecars can rank recent papers by breakthrough-ness the way domain experts do. We GRO-compiled **66 recent (2025–2026) Alzheimer's papers**, rated each with a blind, prior-art-grounded **3-persona expert panel** (0–100), and searched for the best metric — including via a 16→4→2→1 design tournament — under a strict **held-out train/test split**.

Three results, in order of importance:

1. **We are far from the ceiling, and the ceiling is measurable.** A single expert judge predicts the panel at $\rho \approx 0.97$ (leave-one-judge-out) — that is the best any predictor could do. Our best metric reaches $\rho \approx 0.57$ on held-out papers. So the metric captures **~59% of the achievable rank signal**; ~40 points of correlation are lost between *reading the paper* and *the GRO feature vector*. **We have not hit a ceiling** — and the residual gap is *structured* (specific paper genres the features cannot see), not noise, so a better data shape can genuinely close it.
2. **The simplest metric wins; complexity overfits.** The best metric is a **single feature** — the compiler's mean assessed contribution-novelty (`contrib_wmean`), $\rho \approx 0.57$ held-out with *no* train→test gap. Every richer formula the tournament produced climbed the training score and failed to transfer (winner: 0.77 train → 0.53 test). The shipped composite manages only $\rho \approx 0.30$.
3. **The generalizing signal is the compiler's reading of the paper — not the "GRO-native" machinery.** The prior-art anchor and delta-ledger shapes we invested most in did not transfer. This is either a real finding (contribution-typing is what matters) or a warning that our LLM metric and our LLM panel share a bias — distinguishable only on a second corpus. We also tested adding **genre** (fetched deterministically from OpenAlex): it did not improve held-out ranking and is excluded from the metric.
4. **We found a second metric while building the first.** Prior-art *overlap* — how close a paper sits to existing work — is one measurement that reads with **opposite sign for two distinct constructs**: low overlap = **breakthrough/novelty** (distance from prior art), high overlap = **convergence/consolidation** (pulling a field together). We set out to measure the former and, in doing so, defined the latter. Convergence is specified here but not yet validated (it needs its own labeled panel).

The honest deliverable: a **one-feature triage metric** (`contrib_wmean`) at 0.57, a **measured ceiling** at 0.97, a **shared overlap definition** that cleanly separates breakthrough from convergence, and the finding that no cheap structural feature (genre, overlap, delta) beats plain contribution depth on this corpus — so the next move is validation on a second domain, not a richer formula.

1. Features and metrics — definitions (read this first)

Every metric below is deterministic arithmetic over per-paper GRO sidecars. What matters is that the sidecars come from **three sources with different epistemic status** — conflating them is the shipped metric's core error.

Provenance tags used throughout: - **[PAPER→LLM]** — the compiler read *this paper's text* and made a typed judgment. Not externally verified; blind to what happened after publication. - **[PRIOR-ART]** — resolved against *earlier* literature (the paper's references + an OpenAlex search for pre-publication near-neighbors). - **[FOLLOW-ON]** — from *later* papers citing this one (the citation graph).

The features (each defined on first use; full dictionary in `corpus/FEATURES.md`):

feature	provenance	definition
contribution	[PAPER→LLM]	composite $0.6 \cdot \text{peak} + 0.4 \cdot \text{wmean}$ over the paper's typed contributions
contrib_wmean	[PAPER→LLM]	confidence-weighted <i>mean</i> novelty weight across a paper's contributions. Each contribution is typed by the compiler from a closed taxonomy (<code>new_paradigm 1.0 ... replication 0.1</code>); <code>wmean</code> is a <i>mean</i> , so it cannot be inflated by padding a paper with shallow contributions. This is the winning metric.
contrib_peak	[PAPER→LLM]	the single strongest <code>weight × confidence</code> contribution ("does it have ≥ 1 landmark result?")
n_contribs	[PAPER→LLM]	raw count of contributions (breadth; gameable)
n_puffery	[PAPER→LLM]	contributions voided by the anti-puffery lock (claimed but not tied to a real claim)
delta / delta_max / n_deltas	[PAPER→LLM, baseline PRIOR-ART]	the paper's quantified comparisons-vs-prior-work: composite / largest magnitude / count
overlap (min/mean/max)	[PRIOR-ART]	the shared measurement (see SPEC.md): fraction of this paper's contribution already covered by each resolved prior-art neighbor, in [0,1]. <code>mean</code> = neighborhood density; <code>max</code> = closest-predecessor (near-duplicate detector); <code>min</code> = connectedness. Read two ways (§7b): breakthrough = $1 - \text{overlap}$; convergence = $+\text{overlap}$.
resolvability	[PRIOR-ART]	fraction of the paper's references that resolved to an external id — the confidence in the overlap estimate (novelty is discounted when low, so an unresolved neighborhood cannot masquerade as distance). <i>Unreliable in this corpus — see §7b/§8.</i>
anchor	[PRIOR-ART]	legacy composite scoring <i>resolution success</i> — a known-buggy field that rewards that the neighborhood resolved, not its overlap. Superseded by the min/mean/max definition above.
uptake_per_year / cd_index	[FOLLOW-ON]	age-normalized in-corpus citations / disruption index — structurally dead on a 2025–2026 corpus (no downstream edges yet)

Key definition — "contribution." `contributions.yaml` is the compiler's answer to "*what does this paper claim to contribute, and how novel is each claim on its face?*" Each entry is **double-typed**: `author_framed_type` (the authors' pitch) vs `compiler_assessed_type` (the compiler's own call), with a rationale and an anti-puffery lock. So `contrib_wmean` measures **the compiler's reading of the paper's stated contributions, in isolation** — reproducible and auditable, but a judgment, never checked against how the field reacted. That distinction drives every result.

2. Corpus and method

- **Corpus:** 66 recent (2025–2026) Alzheimer's papers in `research/ara-library/`, each GRO-compiled with L8 sidecars; **all 66 SOTA anchors genuinely resolved** (references + live OpenAlex), verified via `corpus/audit_l8.py`.
- **Ground truth — prior-art-grounded expert panel.** Each paper rated 0–100 by three blind personas (clinical trialist, molecular neuroscientist, metascience analyst). Each judge received a **prior-art reading list** — the titles/years of the paper's resolved precedence neighborhood, but **not** the metric's overlap *scores* (to avoid circularity with the anchor feature) — and rated *advance beyond that specific prior art*. Score = **mean** of the three judges (mean, not median, so a lone dissenting judge counts). Raw votes in `corpus/expert_votes_raw_grounded.json`.
- **Held-out evaluation.** 40 train / 26 test, stratified by expert score (`corpus/split.json`). Formula search (including the tournament) happens on **train only**; the test set is scored **once**, by the organizer, via `corpus/apply_formula.py`. No test information ever entered the design loop.
- **Grounding sanity-check.** Re-judging with vs without the prior-art brief agreed at $\rho = 0.97$ on rank but deflated borderline scores (mean 29.6 $\rightarrow$ 24.7) — SOTA context makes judges correctly harsher on work that only *sounds* novel. So the grounded panel is the right target.

3. The shipped metric is weak, for a diagnosable reason

`significance = 0.40·contribution + 0.40·delta + 0.20·anchor` reaches only $\rho \approx \mathbf{0.30}$ (all 66). Per-feature correlation with the grounded panel shows it is worse than its own best part:

feature	ρ vs experts
<code>contribution / contrib_wmean</code>	+0.57 / +0.56 (the real signal)
<code>contrib_peak</code>	+0.46
<code>anchor_n</code>	+0.35
<code>anchor (composite)</code>	+0.20
<code>delta / delta_max / n_deltas</code>	−0.09 / −0.09 / −0.15
<code>anchor_mean_overlap / anchor_max_overlap</code>	−0.39 / −0.49
<code>uptake_per_year / cd_index</code>	−0.03 / −0.08 (dead)

80% of the composite's weight is on `contribution + delta`, and `delta` points the wrong way. Strip to the one good component and it roughly doubles.

4. The tournament, and the overfitting it exposed

A **16→4→2→1 critique-carrying tournament** (`tournament16_workflow.js`): 16 clean-room contestants (Sonnet) each designed a formula, hill-climbing on the train split; a judge critiqued residuals and advanced the strongest; every later round saw the critiques. It converged on a principled-looking winner (`depth × novelty × resolution−confidence`). Then we scored everything on the **held-out test** — the paper's core table:

metric	train (40)	held-out test (26)	all (66)
plain depth $0.5 \cdot \text{peak} + 0.5 \cdot \text{wmean}$	0.546	0.572	0.569
contrib_wmean alone (1 feature)	0.533	0.571	0.564
tournament winner (depth×novelty×conf)	0.771	0.534	0.687
depth × novelty (no gate)	0.693	0.455	0.604
v1 draft (depth – delta report-tells)	0.712	0.449	0.641
novelty alone (1 – max_overlap)	0.627	0.266	0.494
shipped significance	0.141	0.503	0.304

Read the train-vs-test columns: **plain depth has no gap (0.55→0.57)**; every elaboration has a large one. The tournament, by optimizing train ρ , walked into overfitting 40 points. Its real result is *negative*: the added machinery does not generalize.

Honesty note on models (as requested). Contestants ran `claude-sonnet-5` (verified from transcripts). The judges were requested as Fable but **silently fell back to `claude-opus-4-8`** (the `'fable'` alias didn't resolve; it needs `'claude-fable-5'`). Moot for the conclusion — the *selected* winner loses to a one-feature baseline on held-out data regardless of who selected it.

5. The adopted metric

Use `contrib_wmean` (the compiler's confidence-weighted mean assessed contribution-novelty): $\rho \approx 0.57$ vs grounded experts on held-out papers, with no overfitting. Even across persona (trialist / neuroscientist / metascientist ≈ 0.5 – 0.57). It ranks real discoveries on top and floors every annual report and burden re-estimate. A usable **triage prior**, not an oracle. (A rank correlation on $n=26$ has a wide CI, so plain depth does not *significantly* beat the elaborate metrics — but that is the argument: if machinery can't beat one feature on unseen data despite fitting train harder, it isn't justified.)

6. How far from the experts — the ceiling (the headline)

The right yardstick is not zero, it is **how well the experts predict each other**. Leave-one-judge-out — one persona predicting the mean of the other two — gives the best score any predictor could achieve given panel noise:

	ρ
Ceiling (single judge → panel)	0.97
Our best metric (<code>contrib_wmean</code> , held-out)	0.57
Shipped composite	0.30

We are at ~59% of the achievable ceiling — not near it. The ~0.40 gap is the information lost between reading the paper and the GRO feature vector. And it is **structured**, not noise — the biggest metric-vs-expert disagreements cluster on genres the features cannot distinguish:

- `sal26` (plasma-tau biomarker): experts 61, metric floors it — a real advance the compiler under-typed.
- `kes25` (diagnosis review): experts 31, metric ranks it *last* — reviews the field values.

- `cum26` / `wan25c` (drug-pipeline review, burden re-estimate): experts ~6, metric over-rates them — pure genre confusion.

Because the errors are systematic, better data shapes can close the gap. This is not a wall.

7. Why the winner is principled — the sign-flip point

The v1 draft subtracted `delta`. That was theoretically wrong, not just weak. A negatively-correlated feature may be sign-flipped only if it is a **true inverse construct** ($1 - \text{overlap} = \text{novelty}$: causal, generalizes) — not if it is a **confound proxy** (`n_deltas` correlates with breakthrough only because report-genre papers pile up deltas *in this corpus*). Test: *would intervening on the feature causally change breakthrough-ness?* Adding a quantified comparison to a real discovery should not lower its score — so `delta` must not be subtracted. The held-out numbers confirm the theory: the delta-hack underperforms doing nothing.

7b. A second metric fell out — convergence vs breakthrough

The most interesting conceptual result was unplanned. Prior-art **overlap** — computed per resolved neighbor as the fraction of this paper's contribution already covered by that neighbor — is a *single measurement* that a metric can read with **opposite sign for two different constructs**:

- **Breakthrough / novelty** = $1 - \text{overlap}$ (distance from prior work). A genuine advance sits *far* from what exists.
- **Convergence / consolidation** = $+\text{overlap}$ (proximity to, and pulling together of, prior work). A synthesis, a unifying framework, or a field-consolidating review sits *close* to many prior works — and that is valuable, just not *breakthrough*.

We set out to build the first and, in specifying the overlap shape, defined the second. The three aggregates carry distinct meaning and belong in the shared definition:

- **max overlap** — the single closest predecessor. High max = a near-duplicate exists $\Rightarrow$ *not* a breakthrough regardless of the rest (this is what should have caught the `cum26` pipeline review, whose max overlap was 0.93 while its mean looked deceptively novel).
- **mean overlap** — neighborhood density. High mean = a crowded, incremental space; the core convergence signal.
- **min overlap** — connectedness. Very low min = the work is barely attached to any prior art (either genuinely first-in-field, or unresolved — hence the resolvability discount below).

This is a genuine second metric, not a rephrasing. Breakthrough and convergence are different research virtues; a paper can be high on one and low on the other (a bold-but-isolated hypothesis: high breakthrough, low convergence; a masterful synthesis: low breakthrough, high convergence). We validated the breakthrough reading against the expert panel; the **convergence reading is defined but unvalidated** — it needs its own labeled panel ("how much does this paper consolidate the field?"). We flag it as the clearest spin-off hypothesis of this work.

Honest caveat on the current data: neither overlap reading beats plain depth on held-out here, because the 66 per-neighbor overlaps were scored by 66 different resolver agents with no shared rubric. That is a *measurement* problem, not a conceptual one — which is exactly why the shared `overlap` definition (below, and now written into `SPEC.md`) matters before either metric can be trusted.

8. How GRO needs to be extended — to climb toward 0.97

All 16 contestants independently reported the substrate is missing signals the ideal metric needs. The 38 proposals cluster into three asks, and §6's residuals say which would help most:

1. **A standardized overlap definition (min/mean/max) + resolvability** — *now specified in SPEC.md*. The overlap signal is conceptually right (it powers both the breakthrough and convergence readings, §7b) but was unusable here because 66 agents scored it 66 different ways, and `resolvability` was unusable for 38/66 papers (derived from inconsistent compiler flags). The fix is a shared, reproducible overlap measure — ideally computed deterministically from OpenAlex (bibliographic coupling / topic-vector cosine) rather than per-agent LLM judgment — discounted by a resolvability fraction so an unresolved neighborhood cannot masquerade as novelty.

Note on genre (tested, not adopted). Contestants proposed a normalized `genre.primary` field, so we fetched genre deterministically from OpenAlex and tested it. The clean signal available (`type: review`) is redundant with contribution depth, and the surveillance-report distinction is not in OpenAlex `type`; genre-conditioning **hurt** held-out ranking ($0.57 \rightarrow 0.50$). Genre is therefore **excluded from the metric**. A finer, validated genre classifier might help, but it is not the easy win the residuals suggested. 2. **A downstream_assessment / stance edge** (4 proposals) — *highest ceiling, slowest*. Per-citer: does later work *confirm/extend* vs *contest/supersede/fail-to-replicate*? The substrate has citation *counts* (dead) and the compiler's *untested* reading, but nothing for **how the field reacted** — the materialized "cross-object contribution graph" (observation O31), and the only path to measuring *realized* impact and to breaking the shared-LLM-bias worry (§9). 3. **A standalone anchor_resolution_confidence $\in [0,1]$** (14 proposals) — *marginal*. Splits resolution-success from the buggy density reward in `anchor`; the winner had to repurpose `anchor_n` for lack of it.

9. The caveat that outranks the extensions

Our signal (`contrib_wmean`, an LLM's reading) and our target (an LLM panel — inter-judge $\rho \approx 0.96$, implausibly high for independent humans) **may share a model bias**. The 0.57 might be two LLMs agreeing with themselves rather than a transferable notion of breakthrough. We cannot tell on this data. Before building any extension, the cheap decisive test is: **run contrib_wmean against a second, non-AD corpus, ideally with a few human ratings**. If 0.57 holds, the metric is real and the genre/stance extensions are worth building; if it collapses, we have been measuring an echo.

10. Where this leaves the program

The assessment paper said "no metric measures breakthroughs, and we can't compute one on the current shape." This sharpens it: **the shipped composite is weak (0.30); the best available metric is a single contribution-depth feature (0.57 held-out); the achievable ceiling is 0.97; and the ~0.40 gap is fine-grained judgment that no cheap structural feature we tried (genre, overlap, delta) captures**. Enriching the metric with GRO's prior-art/delta machinery overfits and is not recommended. Two things are worth more than a richer formula: (1) validate `contrib_wmean` on a second domain against human raters — to learn whether the 0.57 is real signal or shared LLM bias; (2) with the now-standardized `overlap` definition, build and validate the **convergence** metric (§7b) — the distinct second construct this work surfaced.

Reproduction

All under `research/metrics/v5-breakthrough/`: `corpus/compute_corpus.py` (features $\rightarrow$ `corpus_scored.json`), `corpus/FEATURES.md` (dictionary), `gro_extend_workflow.js` (anchor resolution), `expert_panel_sota_workflow.js` + `corpus/briefs/*` (grounded panel), `corpus/split.json` (40/26), `tournament16_workflow.js` $\rightarrow$ `corpus/tournament16_result.json`, `corpus/apply_formula.py` "`<formula>`" `--split train|test|all` (reproduces every ρ in §4/§6). All correlations recomputed from raw data by the organizer; models read from transcripts (Sonnet contestants, Opus judges).